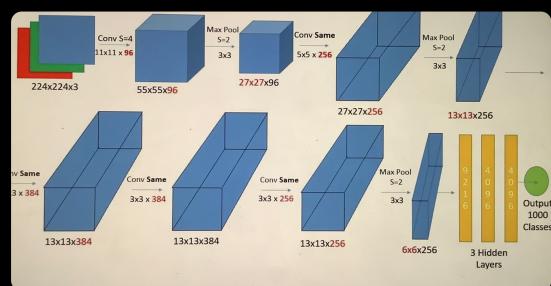


Data Augmentation

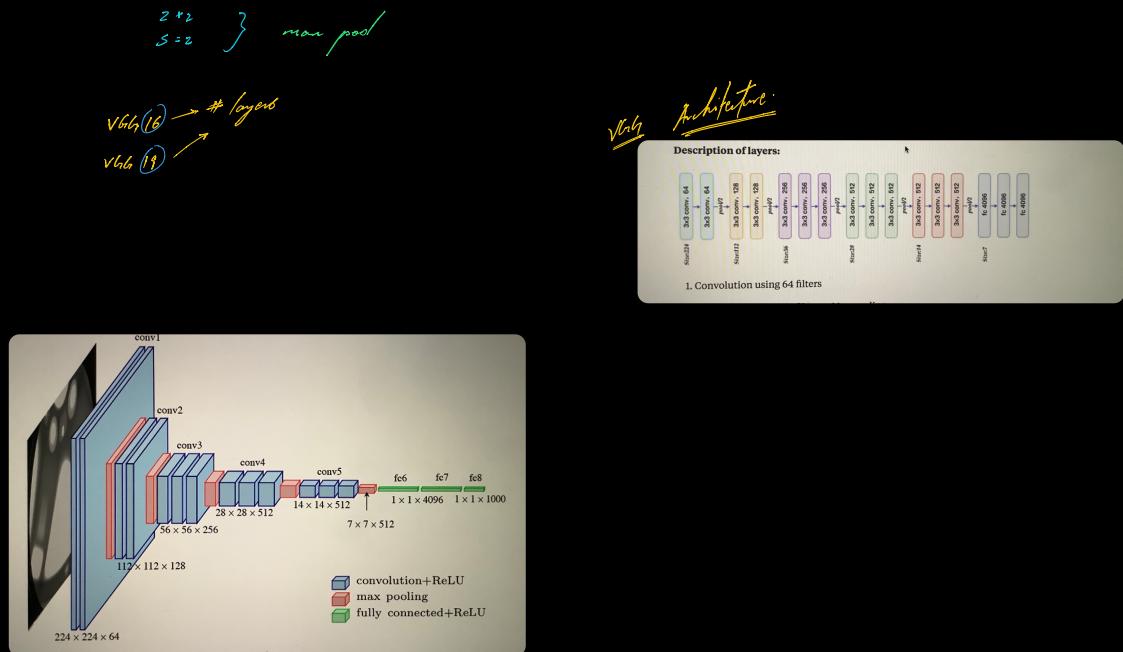


Alex Net



Alex Net
↳ AlexNet++ - better & simplified version

$\begin{matrix} 3 \times 3 \\ s=1 \\ p=\text{same} \end{matrix}$ } of convolutions



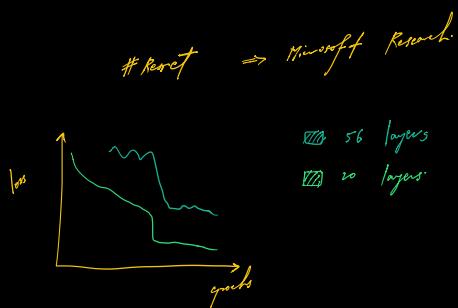
```
# Block 1
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv1')(img_input)
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv2')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block1_pool')(x)

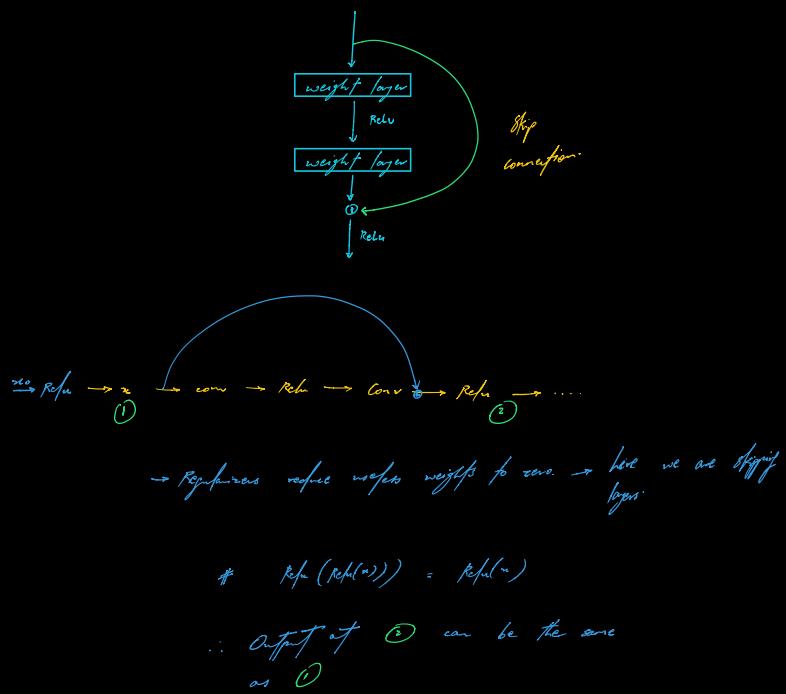
# Block 2
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv1')(x)
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv2')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block2_pool')(x)

# Block 3
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv1')(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv2')(x)
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block3_pool')(x)

# Block 4
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)

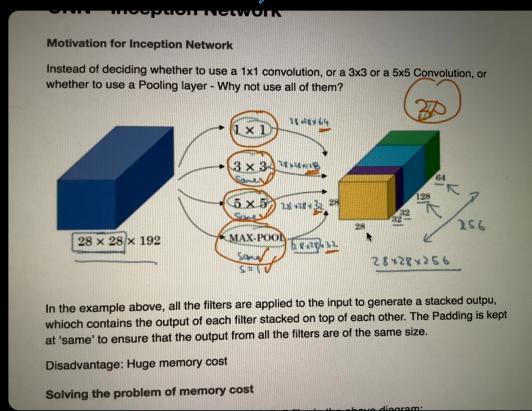
# Block 5
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv3')(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)
```



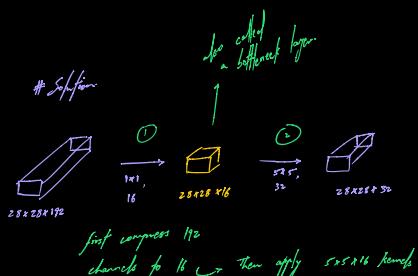
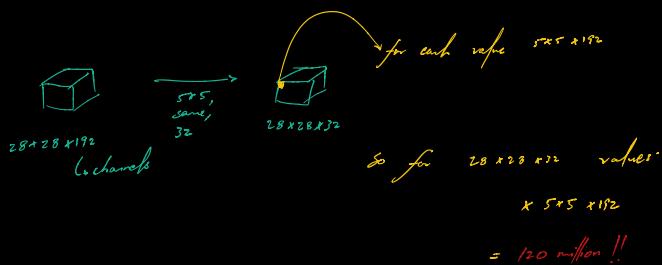


→ This guarantees that applying new layers
 will not hurt → optimizer will make
 weights and layers will be efficient
 → # left off some fucking
 weights than, they made
 float 18 34 50 101 & 152

Inception Net
 1×1 ? 3×3 ? 5×5 ? maxpool?
 Why not have them all!



→ No choosing of kernel size.



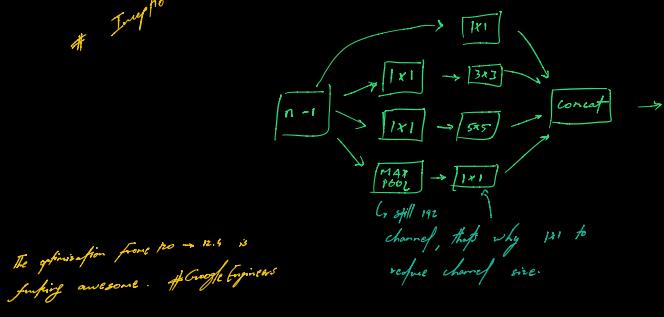
$$\textcircled{1} + \textcircled{2} [28 \times 28 \times 16 \times 16 \times 192] + [28 \times 28 \times 32 \times 5 \times 5 \times 16]$$

$$= 12.4 \text{ million}$$

This is a lot more manageable than 120M

(requested by factor 10!)

Inception Model



Transfer Learning

→ Use pre-trained models instead of noisy random weights.

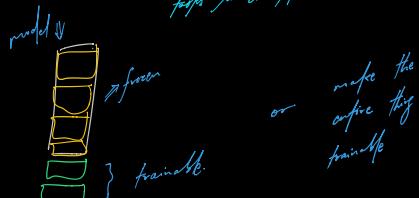
If I'm good at self-building → I'll be good at program in no time

if $[M]$ trained on object A

we can use $[M]$ on object B
given that $A \cup B$ are similar.

TASK A & TASK B
have to be similar

using task will
not be a great
help with analogy,
help to enough.



here role might not
act as a useful
operator for the given
model.

CASE 1 : Task similar to ImageNet but less data

SOLN : Train first few layers, freeze
the rest.

CASE 2 : Similar to ImageNet & large dataset

SOLN : Fine-tune the whole thing \rightarrow small learning rate.

CASE 3 : Small dataset & not similar to ImageNet

SOLN : use output of first couple of
layers as features & flatten.

Use some logic, this thing
depends on the application
we're working on.