

# read the car before the # count

we are going to use VGG  
UpConvolution

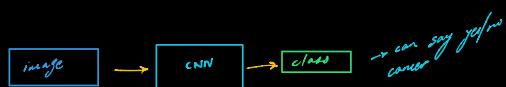


inference of papers on  
image segmentation.

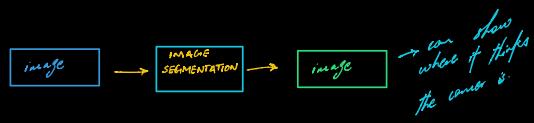
we are going to do:  
the is for called Semantic  
segmentation

never form

not a raw output all  
meaningful segmentation



# background blur  
# soft driving car  
# medical - mark the tumor/ problems → CT scans, etc.  
# colorization

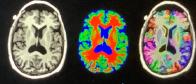


# something like this:

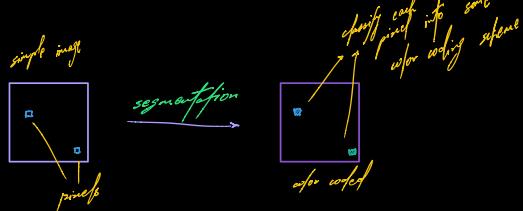


→ detect where  
the nerves are  
in retina scan

#MRI



→ biomedical imaging  
↳ first application of VGG



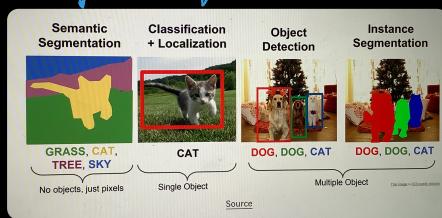
group dataset  
that are general purpose  
datasets

# coco input → ImageNet Equivalent for  
segmentation:  
↳ 1.5 million mapped/fixed label  
↳ 330k images  
↳ 80 classes  
example

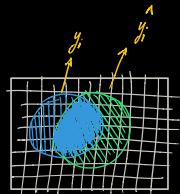


domain  
 sparsity  
 dataset  
 # enhance → for traffic & self driving cars.  
 # traffic light segmentation

# segmentation is object detection → faster & easier  
 or less pixels  
 # segmentation is just object detection



segmentation is just object detection  
 but pixel level, instead of bouding  
 box.



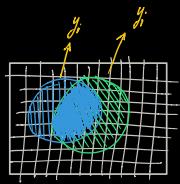
### # Evaluation Metric

○ Average precision & recall for each class → pixel level.

$$\textcircled{c} \text{ Jaccard Similarity} = \frac{R_i \cap R_i^*}{R_i \cup R_i^*}$$

Intersection  
 Over  
 Union  
 $|I \cup U|$

$R$  stands for region



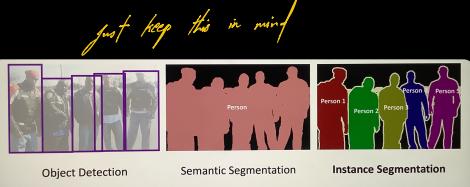
Ideal case =  $I \rightarrow \cap \& \cup$  is same.  
 worst case =  $\cap \approx 0 \rightarrow$  prediction too far.

### # Classical Approaches

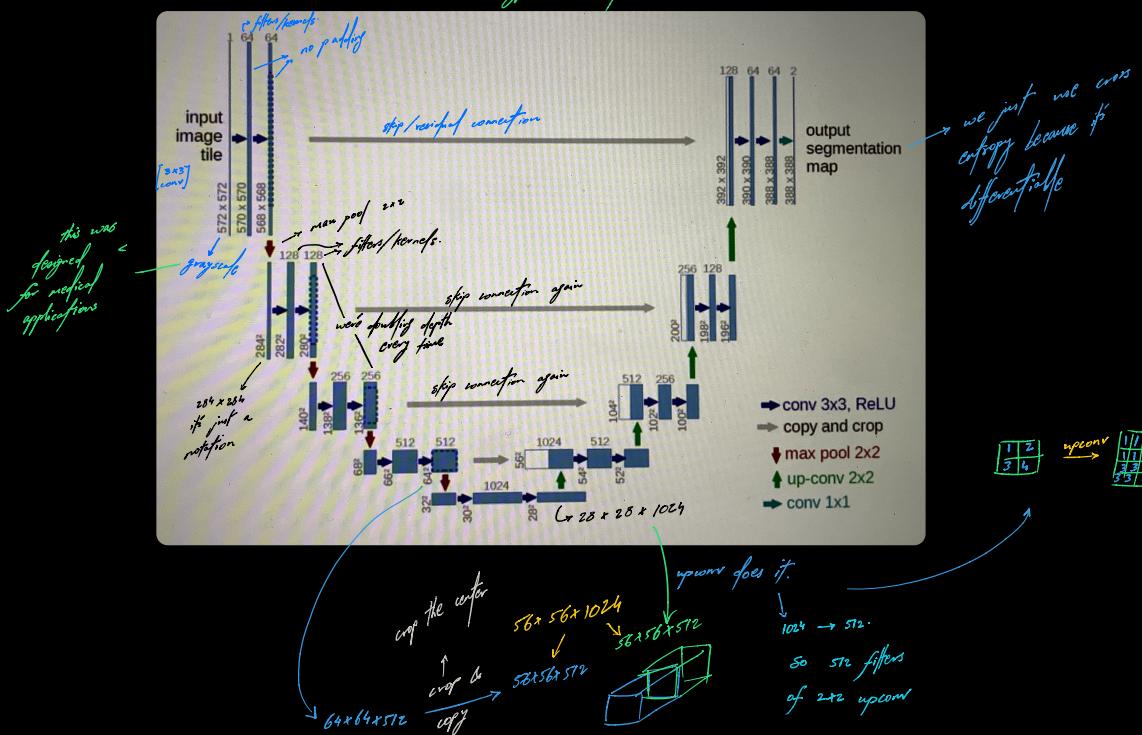
# Clustering using k-means etc → easy to mess up with cluttered images  
 # Edge Detection → traditional CV methods  
 # Graph Theory → they were state of the art, but no more. RIP.

### # Deep Learning Approaches

- CHANCE FOR THIS SESSION
- # UNet - made for biomedical but works everywhere
  - # Fully Convolutional Networks (FCN)
  - # Mask R-CNN
  - # SegNet → for mobile devices



## Encoder + Decoder type architecture



```

def unet(pretrained_weights = None, input_size = (256, 256, 1)):
    inputs = Input(input_size)
    conv1 = Conv2D(64, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(inputs)
    conv1 = Conv2D(64, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(conv1)
    pool1 = MaxPooling2D(pool_size=(2, 2))(conv1)
    conv2 = Conv2D(128, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(pool1)
    conv2 = Conv2D(128, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(conv2)
    pool2 = MaxPooling2D(pool_size=(2, 2))(conv2)
    conv3 = Conv2D(256, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(pool2)
    conv3 = Conv2D(256, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(conv3)
    pool3 = MaxPooling2D(pool_size=(2, 2))(conv3)
    conv4 = Conv2D(512, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(pool3)
    conv4 = Conv2D(512, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(conv4)
    drop4 = Dropout(0.5)(conv4) → forget was not here but this code has it.
    pool4 = MaxPooling2D(pool_size=(2, 2))(drop4)

    conv5 = Conv2D(1024, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(pool4)
    conv5 = Conv2D(1024, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(conv5)
    drop5 = Dropout(0.5)(conv5)

→ after axis 3 go from 1024 to 512
    up6 = Conv2D(512, 2, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(UpSampling2D(size = (2,2))(drop5))
    merge6 = concatenate([drop4, up6], axis = 3)
    conv6 = Conv2D(512, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(merge6)
    conv6 = Conv2D(512, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(conv6)

    up7 = Conv2D(256, 2, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(UpSampling2D(size = (2,2))(conv6))
    merge7 = concatenate([conv3, up7], axis = 3)
    conv7 = Conv2D(256, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(merge7)
    conv7 = Conv2D(256, 3, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(conv7)

    up8 = Conv2D(128, 2, activation = 'relu', padding = 'same', kernel_initializer = 'he_normal')(UpSampling2D(size = (2,2))(conv7))
    merge8 = concatenate([conv2, up8], axis = 3)

```

*# This is fully  
singular due to  
dropout!*

*# they haven't  
cropped the top  
convolutions because  
the dimensions  
line up*

*# Data Augmentation is necessary  
→ shift, rotate, shear, zoom, etc (cropper graphic) → some applications  
won't allow this.  
→ grayscale adjustments*

*# resources  
# Harrell Lamb's flag  
# Google N's flag on this  
# Divamit's repo on GitHub for  
popular implementations.*

*mark region out  
→ does bounding box  
→ then highlights object point wise*

*This is just the architecture,  
we can change the %o exception  
as long as the numbers add up.*

*follow top researchers  
in conferences  
→ that's how we filter → finding good papers  
or Arxiv Sanity  
one good paper a week  
the more*