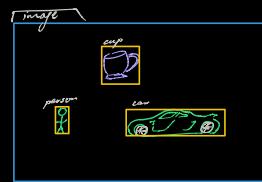


go through image segmentation first
to get a better idea of the scenario

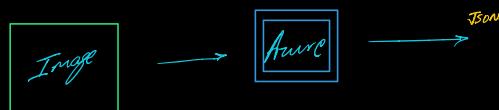
Well see you vs → there's another senior for this
↓
written in C
↳ # good

- * companies have performed models
and they're available as APIs.
In this session we will use Amazon API



have 10% foamt require credit card
to by. google does require it

* They have lots of features in the
MS a excellent documentation of
all of that. I don't care, so
I haven't done. we can just look it up for any
info



Limitations

It's important to note the limitations of object detection so you can avoid or mitigate the effects of false negatives (missed objects) and limited detail.

- Objects are generally not detected if they're small (less than 5% of the image).
 - Objects are generally not detected if they're arranged closely together (a stack of plates, for example).
 - Objects are not differentiated by brand or product names (different types of sodas on a store shelf, for example). However, you can get brand information from an image by using the [Brand detection](#) feature.

you, etc all have good fun

```
{
  "objects": [
    {
      "rectangle": {
        "x": 730,
        "y": 66,
        "w": 135,
        "h": 85
      },
      "object": "kitchen appliance",
      "confidence": 0.501
    },
    {
      "rectangle": {
        "x": 523,
        "y": 377,
        "w": 185,
        "h": 46
      },
      "object": "coffee cup",
      "confidence": 0.46
    },
    {
      "rectangle": {
        "x": 523,
        "y": 218,
        "w": 189,
        "h": 226
      },
      "object": "computer keyboard",
      "confidence": 0.51
    },
    {
      "rectangle": {
        "x": 473,
        "y": 289,
        "w": 189,
        "h": 226
      },
      "object": "Laptop",
      "confidence": 0.85,
      "parent": {
        "object": "computer",
        "confidence": 0.851
      }
    }
  ]
}
```

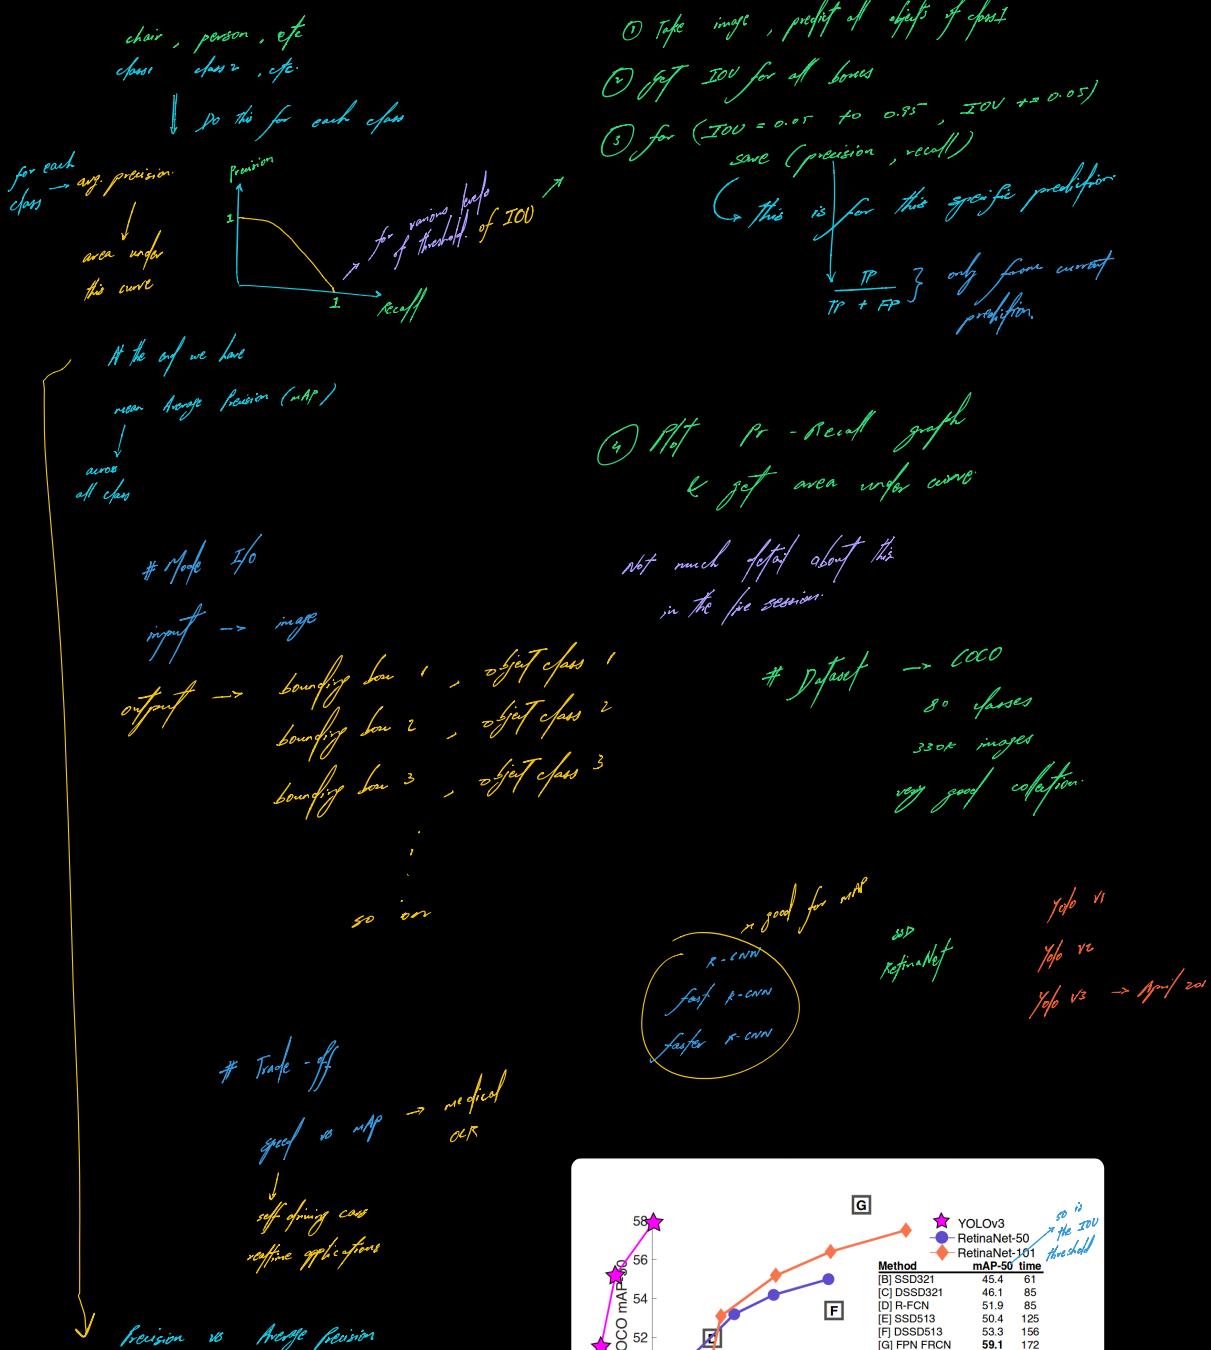
Performance Refin

- # if we print json simply,
it will be non-injected
- # this function prints json clearly
print(`json.dumps(json_file)`)

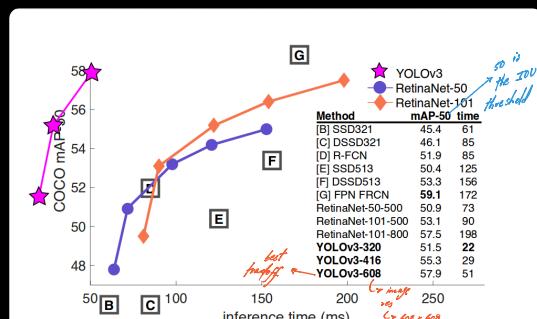
$$IoU = \frac{\text{Area of Intersection}}{\text{Area of Union}}$$

↓
intersection over Union

→ this is just for one
box.



Precision refers to precision at a particular decision threshold. For example, if you count any model output less than 0.5 as negative, and greater than 0.5 as positive. But sometimes (especially if your classes are not balanced, or if you want to favor precision over recall or vice versa), you may want to vary this threshold. Average precision gives you average precision at all such possible thresholds, which is also similar to the area under the precision-recall curve. It is a useful metric to compare how well models are ordering the predictions, without considering any specific decision threshold.



YOLO vs - feature extractor (Darknet-53)

fully convolutional network (FCN) → no maxpool

conv → bottleneck → bottleneck

no maxpooling → they use stride = 2 for conv
kind works the same.

pretrained on ImageNet.

skip connections allow deeper → deeper architectures.

↓
YOLO vs net
Darknet53 (+ depth have
skip connections)

Type	Filters	Size	Output
Convolutional	32	3×3	256 × 256
Convolutional	64	$3 \times 3 / 2$	128 × 128
Convolutional	32	1×1	
1x1 Convolutional	64	3×3	128 × 128
Residual			
Convolutional	128	$3 \times 3 / 2$	64 × 64
Convolutional	64	1×1	
Convolutional	128	3×3	
Residual			64 × 64
Convolutional	256	$3 \times 3 / 2$	32 × 32
Convolutional	128	1×1	
Convolutional	256	3×3	
Residual			32 × 32
Convolutional	512	$3 \times 3 / 2$	16 × 16
Convolutional	256	1×1	
Convolutional	512	3×3	
Residual			16 × 16
Convolutional	1024	$3 \times 3 / 2$	8 × 8
Convolutional	512	1×1	
4x Convolutional	1024	3×3	
Residual			8 × 8
Avgpool		Global	
Connected		1000	
Softmax			

Darknet-53 model

53 conv layers

why not use ResNet as
feature extractor? It generally has better performance.

It has same Top1, Top5 but half
the FLOPs. Performance surprise. That's why.

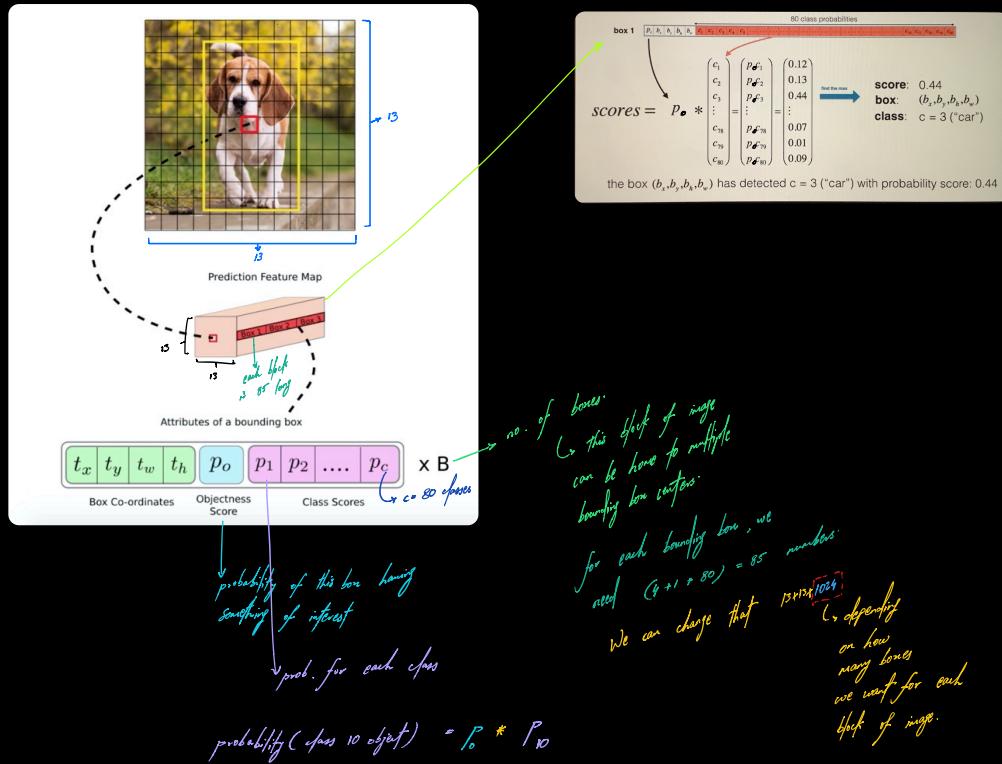
↓
YOLO v3 → 78 FPS
YOLO v2 → 171 FPS
Darknet53 → 37 FPS → MEGH

well, we YOLOv3 is an
example, it's balanced yield + accuracy.

input image → $448 \times 448 \times 3$
feature extracted → $13 \times 13 \times 1024$
↳ learned of conv stride = 2 → 5 of them
↳ 1024 filters in the last layer

↓
from Darknet53
pretrained on ImageNet

* Drawing boxes & obj



↳ no. of boxes
↳ this block of image
can be home to multiple
bounding box centers
for each bounding box, we
need $(4 + 80) = 85$ numbers
↳ we can change that $13 \times 13 \times 1024$
↳ depending
on how
many boxes
we want for each
block of image.

for eg. if I want 5 po
for each block of image

$$13 \times 13 \times (4 \times (5 + 80))$$

no. of boxes = 5 classes = 80

$$\therefore 13 \times 13 \times 420$$

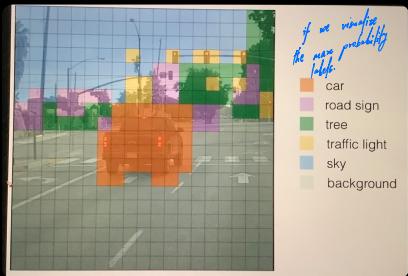
in yolo they use 13×13
so they much less effort to predict
smaller $B = 5$
larger $B = 10$ will

all this is differentiable, so it works

that gives us $13 \times 13 \times 1024$ the ...

that are often convolutions involved
 $\therefore 13 \times 13 \times 405 \rightarrow$ they need
this is simplified version

p_1 conv
they need the thing to
optimize function

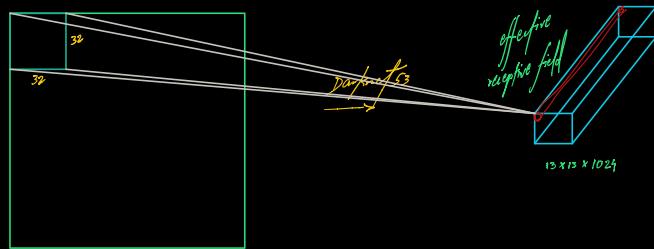


→ intermediate result
of 100×100

* If there's only one object, then P^o

↓
probability of
having an object
of interest will be
zero or close to
zero

* YOLO is designed to work
on square images but
we can always resize
the input to be a square

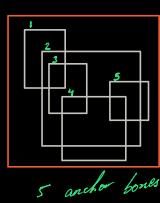


* what if the object
spans multiple receptive fields?

↓
the output of NN has
to work receptive field, the
height weight can be anything

↓
anchor boxes solve
some of the issue

↑
 $\text{range} \rightarrow 0^{+0.2}$ → it'll be at most one
pixel away. Try not
bb to be very close to
another border



* solution → have anchor boxes

* left represent bb as a modification
of anchor boxes

$$(\text{cx}, \text{cy}, \text{pw}, \text{ph})$$

The improved performance by 3x or 4x
by using anchor boxes

sigmoid function

$$\sigma(cx) + cx$$

$$\sigma(cy) + cy$$

$$(\text{bx}, \text{by}, \text{bw}, \text{bh})$$

$$\begin{cases} \text{bx} \\ \text{by} \\ \text{bw} \\ \text{bh} \end{cases} = \begin{cases} \sigma(cx) + cx \\ \sigma(cy) + cy \\ \text{pw} \cdot e^{cw} \\ \text{ph} \cdot e^{ch} \end{cases}$$